SandboxOps

Blast Radius Zero for AI agents

plan → sandbox → verify → evidence

Igor Eduardo · Vultr Agent Rush (lablab.ai)

The problem

Agents that write and run code are useful.
Giving them a shell on your server is not safe.

Ops teams need to know **what ran** and that it **could not do damage**.

The containment moment

Agent-written code tries to call the internet:

```
socket.create_connection(('example.com',80))
```

🛑 CONTAINED: network disabled in sandbox

Infinite loop → killed by CPU limit

How it works

1. Plan: natural language → Python (Vultr Serverless Inference)
2. Sandbox: throwaway container, no network, read-only, non-root, caps dropped
3. Verify: output checked against a rule
4. Evidence: code + output + exit code + SHA-256

Architecture

Browser → SandboxOps app (Vultr Cloud Compute VM, port 80, /health)
→ Vultr Serverless Inference (api.vultrinference.com/v1)
→ docker run --network none --read-only --cap-drop ALL
→ verify → evidence record

Vultr usage

- Vultr Cloud Compute VM: web app + Docker sandbox host (systemd)
- Vultr Serverless Inference: the planner model (OpenAI-compatible API)
- One-command deploy: `deploy/vultr-deploy.sh`

Results

- test_agent.py: 3/3 (bill verified at $551.00, network blocked, loop killed)
- test_app.py: 5/5 (/health, verified run, 2 containment runs, UI)
- Open source, MIT

Roadmap & buyer

Buyer: finance and healthcare ops teams running LLM-written reconciliations and bill checks

Next: gVisor runtime · egress allowlist policy · signed evidence · audit export

github.com/igoreduardoai/sandboxops